

# Build Your First RAG Chatbot – Complete Beginner Guide

---

## What You'll Learn

By the end of this tutorial, you'll understand:

- What RAG (Retrieval Augmented Generation) means
- How to build a chatbot that answers from your documents
- Vector databases and embeddings (simplified!)
- How to prevent AI hallucination

## Prerequisites

- Basic Python knowledge (variables, functions, loops)
- Know how to run Python scripts
- Understand what APIs are

## Key Libraries & Tools We'll Use



### Google Generative AI (genai)

- **What it does:** The "brain" that generates human-like responses
- **Why we use it:** Fast, accurate, and follows instructions well
- **Think of it as:** A very smart assistant that can read and write



### Sentence Transformers

- **What it does:** Converts text into numbers (vectors) that computers can compare
- **Why we use it:** Finds similar meanings even with different words
- **Think of it as:** A translator that turns words into math
- **Example:** "cat" and "kitten" get similar numbers because they're related



### ChromaDB

- **What it does:** Stores and searches through millions of number vectors super fast
- **Why we use it:** Like Google search but for your document chunks
- **Think of it as:** A super-smart filing cabinet that finds related documents instantly
- **Cool feature:** Remembers everything even after you restart the program

## PyPDF2

- **What it does:** Reads text from PDF files
- **Why we use it:** Most documents are in PDF format
- **Think of it as:** A tool that copies text from PDFs so our AI can read it

## Streamlit

- **What it does:** Creates web apps with just Python (no HTML/CSS needed!)
- **Why we use it:** Makes a beautiful chat interface in minutes
- **Think of it as:** Magic that turns Python scripts into websites

## Python-dotenv

- **What it does:** Safely stores secret keys (like API keys)
- **Why we use it:** Keeps your API keys secure and separate from code
- **Think of it as:** A secure vault for passwords

## Key Terms You'll Learn

**Vector/Embedding:** A list of numbers that represents the "meaning" of text

`"I love cats" → [0.2, 0.8, 0.1, 0.9, ...] # 384 numbers!`

**Similarity Search:** Finding text with similar meanings by comparing their numbers

`"What are cats?" finds "Cats are pets" because their numbers are similar`

**Chunking:** Breaking long documents into smaller pieces

`"Long document..." → ["Piece 1", "Piece 2", "Piece 3"]`

**Context:** The relevant document pieces we show to the AI

```
User asks: "What do cats eat?"
```

```
Context: "Cats eat fish. Cats like milk. Cats hunt mice."
```

```
AI answers using only this context
```

**RAG Pipeline:** The complete process

```
Documents → Chunks → Vectors → Database → Search → Context → AI →  
Answer
```

## Step-by-Step Learning Path

### Phase 1: Understanding the Concepts (15 minutes)

#### What is RAG?

Think of RAG like a smart student taking an open-book exam:

- **Regular AI:** Answers from memory (might make mistakes)
- **RAG AI:** Looks up answers in provided books first, then responds

#### The Magic Behind RAG

```
Your Documents → Convert to Numbers → Store in Database → Search →  
Answer
```

1. **Documents:** Your PDFs, text files
2. **Convert to Numbers:** Each piece of text becomes a list of numbers (vectors)
3. **Store:** Save these numbers in a searchable database
4. **Search:** Find similar numbers when you ask a question
5. **Answer:** AI reads only the found text and responds






### Phase 2: Setting Up Your Environment (10 minutes)

#### Install Python Tools

```
# Create a new folder for your project  
mkdir my-rag-chatbot  
cd my-rag-chatbot
```

```
# Install required packages (this might take a few minutes)
pip install streamlit sentence-transformers chromadb google-generativeai
```

## What each package does:

- `streamlit` → Web interface magic ✨
- `sentence-transformers` → Text-to-numbers converter 
- `chromadb` → Super-fast search database 
- `google-generativeai` → The AI brain 
- `python-dotenv` → Secret key manager 
- `pypdf2` → PDF text extractor 



## Get Your AI API Key

1. Go to [Google AI Studio](#)
2. Create a free account
3. Generate an API key
4. Save it somewhere safe

## Phase 3: Build Step by Step

### Step 1: Create the Brain (15 minutes)

Create `simple_rag.py`:

```
# This is like the brain of our chatbot
import google.generativeai as genai # 🧠 Google's AI
from sentence_transformers import SentenceTransformer #  Text → Number
import chromadb #  Vector database

class SimpleChatbot:
    def __init__(self, api_key):
        # Set up the AI (Google Gemini)
        genai.configure(api_key=api_key)
        self.ai = genai.GenerativeModel('gemini-pro')

        # Set up the "memory" (converts text to numbers)
        # This model is small, fast, and good quality
        self.memory = SentenceTransformer('all-MiniLM-L6-v2')

        # Set up the "filing cabinet" (stores the numbers)
        self.cabinet = chromadb.Client()
```

```

self.files = self.cabinet.create_collection("my_docs")

def learn_document(self, text, doc_name):
    """Teach the chatbot about a document"""
    # Break text into small pieces
    pieces = [text[i:i+500] for i in range(0, len(text), 400)]

    # Convert pieces to numbers
    numbers = self.memory.encode(pieces)

    # Store in filing cabinet
    ids = [f"{doc_name}_{i}" for i in range(len(pieces))]
    self.files.add(
        embeddings=numbers.tolist(),
        documents=pieces,
        ids=ids
    )
    print(f"✅ Learned about {doc_name}")

def answer_question(self, question):
    """Answer a question using learned documents"""
    # Convert question to numbers
    question_numbers = self.memory.encode([question])

    # Search filing cabinet for similar pieces
    results = self.files.query(
        query_embeddings=question_numbers.tolist(),
        n_results=3
    )

    # Combine found pieces
    context = "\n".join(results['documents'][0])

    # Ask AI to answer using only these pieces
    prompt = f"""
    Answer this question using ONLY the provided context.
    If the answer isn't in the context, say "I don't know."

    Context: {context}
    Question: {question}
    Answer:
    """

    response = self.ai.generate_content(prompt)

```

```

        return response.text

# Test it!
if __name__ == "__main__":
    # Replace with your API key
    bot = SimpleChatbot("your-api-key-here")

    # Teach it something
    bot.learn_document("Python is a programming language. It's easy to le

    # Ask it something
    answer = bot.answer_question("What is Python?")
    print(answer)


```

**Try this first!** Run it and see if it works.

## Step 2: Add File Reading (10 minutes)

Create `file_reader.py`:

```

import PyPDF2 #  The PDF text extractor

def read_pdf(file_path):
    """Read text from a PDF file using PyPDF2"""
    with open(file_path, 'rb') as file: # 'rb' = read binary
        reader = PyPDF2.PdfReader(file) # Create PDF reader
        text = ""
        # Loop through each page and extract text
        for page in reader.pages:
            text += page.extract_text()
    return text

def read_txt(file_path):
    """Read text from a TXT file"""
    with open(file_path, 'r') as file:
        return file.read()

# Test it!
if __name__ == "__main__":
    # Create a test file
    with open("test.txt", "w") as f:
        f.write("This is a test document about cats. Cats are amazing pet

```

```
text = read_txt("test.txt")
print(f"Read: {text}")
```

### Step 3: Combine Everything (15 minutes)

Create `complete_chatbot.py`:

```
from simple_rag import SimpleChatbot
from file_reader import read_pdf, read_txt
import os

def main():
    # Get API key
    api_key = input("Enter your Google AI API key: ")

    # Create chatbot
    bot = SimpleChatbot(api_key)

    # Load documents from a folder
    docs_folder = "documents"
    if not os.path.exists(docs_folder):
        os.makedirs(docs_folder)
        print(f"Created {docs_folder} folder. Add your PDF/TXT files there")
        return

    # Read all files in the folder
    for filename in os.listdir(docs_folder):
        file_path = os.path.join(docs_folder, filename)

        if filename.endswith('.pdf'):
            text = read_pdf(file_path)
            bot.learn_document(text, filename)
        elif filename.endswith('.txt'):
            text = read_txt(file_path)
            bot.learn_document(text, filename)

    # Chat loop
    print("\n🤖 Chatbot ready! Ask me about your documents (type 'quit' to exit)")
    while True:
        question = input("\nYou: ")
        if question.lower() == 'quit':
            break
```

```

        answer = bot.answer_question(question)
        print(f"Bot: {answer}")

if __name__ == "__main__":
    main()

```

## Step 4: Add Web Interface (20 minutes)

Create `web_app.py`:

```

import streamlit as st
from complete_chatbot import SimpleChatbot
import os

st.title("🤖 My First RAG Chatbot")

# Get API key
api_key = st.text_input("Enter your Google AI API key:", type="password")

if api_key:
    # Create chatbot
    if 'bot' not in st.session_state:
        st.session_state.bot = SimpleChatbot(api_key)

    # Load documents
    docs_folder = "documents"
    if os.path.exists(docs_folder):
        for filename in os.listdir(docs_folder):
            if filename.endswith(('.pdf', '.txt')):
                # Load document logic here
                st.success(f"Loaded {filename}")

    # Chat interface
    question = st.text_input("Ask me about your documents:")

    if question:
        answer = st.session_state.bot.answer_question(question)
        st.write(f"**Answer:** {answer}")

```

## Phase 4: Understanding What You Built

### The Vector Magic Explained



```
# When you add "Cats are pets"
text = "Cats are pets"
vector = [0.1, 0.5, 0.8, 0.2, ...] # 384 numbers!

# When you ask "What are cats?"
question = "What are cats?"
q_vector = [0.2, 0.4, 0.9, 0.1, ...] # Similar numbers!

# The system finds similar vectors and returns the text
```

## Why This Prevents Hallucination

- AI only sees the text pieces you found
- No access to its training data
- Must say "I don't know" if answer isn't in your documents

## Phase 5: Common Beginner Mistakes & Solutions

### Mistake 1: "My chatbot gives weird answers"

**Problem:** Chunks are too small or too big

**Solution:** Adjust chunk size (try 500-1000 characters)

### Mistake 2: "It says 'I don't know' for everything"

**Problem:** Search isn't finding relevant pieces

**Solution:** Check if documents loaded correctly, try different search terms

### Mistake 3: "It's too slow"

**Problem:** Processing large documents

**Solution:** Use smaller documents for testing first

## Phase 6: Next Steps

### Make It Better

1. **Add more file types:** Word docs, web pages
2. **Better chunking:** Split by sentences, not characters
3. **Multiple languages:** Use different embedding models
4. **Better UI:** Add file upload, chat history
5. **Deploy it:** Put it online with Heroku or Streamlit Cloud

## Advanced Concepts to Learn Later

- **Hybrid search:** Combine keyword + vector search
- **Reranking:** Improve search results
- **Fine-tuning:** Train models on your specific data
- **Evaluation:** Measure how good your answers are

## Quick Reference

### Essential Commands

```
# Install everything
pip install streamlit sentence-transformers chromadb google-generativeai

# Run web app
streamlit run web_app.py

# Run command line version
python complete_chatbot.py
```

## Key Concepts & Libraries Summary

### Core Libraries

- **google-generativeai** : Google's AI that generates responses
- **sentence-transformers** : Converts text to searchable numbers
- **chromadb** : Fast vector database for storing document chunks
- **pypdf2** : Reads text from PDF files
- **streamlit** : Creates web interfaces with Python
- **python-dotenv** : Manages environment variables securely

### Essential Concepts

- **Embedding/Vector:** Text converted to numbers for similarity comparison
- **Vector Database:** Storage system optimized for finding similar vectors
- **Similarity Search:** Finding related content by comparing number patterns
- **Context Window:** Amount of text the AI can process at once
- **Chunking:** Splitting documents into manageable pieces
- **RAG Pipeline:** Document → Vector → Search → Context → AI Response

# Troubleshooting

## "ModuleNotFoundError"

```
pip install [missing-module-name]
```

## "API Key Error"

- Check your Google AI Studio account
- Make sure you copied the key correctly
- Try generating a new key

## "No documents found"

- Make sure files are in the `documents/` folder
- Check file extensions (.pdf, .txt)
- Try with a simple .txt file first

## Congratulations! 🎉

You've built a RAG chatbot that:

- ✅ Reads your documents
- ✅ Converts them to searchable vectors
- ✅ Finds relevant information
- ✅ Answers questions without hallucination
- ✅ Has a web interface

You now understand the core concepts behind modern AI applications!